

# CSCI 2021, Spring 2016

## Optimizing the Performance of a Pipelined Processor

Assigned: March 28, Due: April 13, 11:55PM

Nilanshu Sharma ([sharm398@umn.edu](mailto:sharm398@umn.edu)) is the lead person for this assignment.

### 1 Introduction

In this lab, you will learn about the design and implementation of a pipelined Y86-64 processor, optimizing both it and a benchmark program to maximize performance. You are allowed to make any semantics-preserving transformation to the benchmark program, or to make enhancements to the pipelined processor, or both. When you have completed the lab, you will have a keen appreciation for the interactions between code and hardware that affect the performance of your programs.

The lab is organized into three parts, each with its own handin. In Part A you will write some simple Y86-64 programs and become familiar with the Y86-64 tools. In Part B, you will extend the SEQ simulator with a new instruction. These two parts will prepare you for Part C, the heart of the lab, where you will optimize the Y86-64 benchmark program and the processor design.

### 2 Logistics

You will work on this lab alone.

Any clarifications and revisions to the assignment will be posted on the course Web page.

### 3 Handout Instructions

1. Start by copying the file `archlab-handout.tar` to a (protected) directory in which you plan to do your work.
2. Then give the command: `tar xvf archlab-handout.tar`. This will cause the following files to be unpacked into the directory: `README`, `Makefile`, `sim.tar`, `archlab.pdf`, and `simguide.pdf`.

3. Next, give the command `tar xvf sim.tar`. This will create the directory `sim`, which contains your personal copy of the Y86-64 tools. You will be doing all of your work inside this directory.
4. Finally, change to the `sim` directory and build the Y86-64 tools:

```
unix> cd sim
unix> make clean; make
```

Note: If you're working on cselab machines and `make` fails with the following error:

```
/usr/bin/ld: cannot find -ltk
/usr/bin/ld: cannot find -ltcl
collect2: ld returned 1 exit status
make[1]: *** [psim] Error 1
```

Make the following changes to Makefile in `sim` directory

```
[Old] -> TKLIBS=-L/usr/lib -ltk -ltcl
```

```
[New] -> TKLIBS=-L/usr/lib -ltk8.5 -ltcl8.5
```

If you are working on your local machine, put the appropriate version of `ltk` and `ltcl` in the Makefile.

If the you see this statement in the Makefile:

```
#GUIMODE=-DHAS_GUI
```

Uncomment it.

## 4 Part A

You will be working in directory `sim/misc` in this part.

Your task is to write and simulate the following two Y86-64 programs. The required behavior of these programs is defined by the example C functions in `examples.c`. Be sure to put your name and ID in a comment at the beginning of each program. You can test your programs by first assembling them with the program `YAS` and then running them with the instruction set simulator `YIS`.

In all of your Y86-64 functions, you should follow the x86-64 conventions for passing function arguments, using registers, and using the stack. This includes saving and restoring any callee-save registers that you use.

### **max.y86: Iteratively find the max in a linked list**

Write a Y86-64 program `max.y86` that iteratively finds the max among elements of a linked list. Your program should consist of some code that sets up the stack structure, invokes a function, and then halts. In this case, the function should be Y86-64 code for a function (`max_list`) that is functionally equivalent to the C `max_list` function in Figure 1. Test your program using the following three-element list:

```
# Sample linked list
.align 8
ele1:
    .quad 0x00a
    .quad ele2
ele2:
    .quad 0x0b0
    .quad ele3
ele3:
    .quad 0xc00
    .quad 0
```

### **matrix\_and\_xor.y86: Perform a matrix AND**

Write a Y86 program (`matrix_and_xor.y86`) that performs a sequence of Bitwise AND operations on corresponding elements two matrices and stores the result into another matrix. Your `matrix_and_xor` function should take two matrices `A[i][j]` and `B[i][j]`. The result matrix (`C[i][j]`) is the bitwise AND of `A[i][j]` and `B[i][j]`. It is similar to matrix addition/subtraction function with bitwise AND instead of addition/subtraction. After calculating the matrix `C`, take the bitwise XOR of every element in the matrix `C` and store the result in register `%rax`. Test your program using the following array representation of matrix. Consider the matrices as having 2x2 dimensions. You will be given 2 arrays each of size 4, one corresponding to matrix `A` and other corresponding to matrix `B`. Element 0 of the following array is interpreted as `mat[0][0]`. Similarly Element 1 is interpreted as `mat[0][1]`, Element 2 as `mat[1][0]` and Element 4 as `mat[1][1]`.

```
# Matrix A
    .align 8

matA:
    .quad 0x001
    .quad 0x002
    .quad 0x003
    .quad 0x004

# Matrix B

matB:
    .quad 0x005
    .quad 0x006
    .quad 0x007
```

```

1 /* linked list element */
2 typedef struct ELE {
3     long val;
4     struct ELE *next;
5 } *list_ptr;
6
7 /* max_list - Find the maximum element in a linked list */
8 long max_list(list_ptr ls)
9 {
10     long max_element = 0;
11     while (ls) {
12         if(ls->val > max_element) {
13             max_element = ls->val;
14         }
15         ls = ls->next;
16     }
17     return max_element;
18 }
19
20 /* matrix_and_xor - Bitwise AND corresponding elements of two matrices and take XOR*/
21 long matrix_and_xor(long size, long A[size][size], long B[size][size],
22                     long C[size][size])
23 {
24     long i,j, result=0;
25     for(i=0; i<size; i++) {
26         for(j=0; j<size; j++) {
27             C[i][j] = A[i][j] & B[i][j];
28         }
29     }
30     for(i=0; i<size; i++) {
31         for(j=0; j<size; j++) {
32             result ^= C[i][j];
33         }
34     }
35     return result;
36 }
37
38 /* copy_block - Copy src to dest and return xor checksum of src */
39 long copy_block(long *src, long *dest, long len)
40 {
41     long result = 0;
42     while (len > 0) {
43         long val = *src++;
44         *dest++ = val;
45         result ^= val;
46         len--;
47     }
48     return result;
49 }
50 }

```

Figure 1: **C versions of the Y86-64 solution functions.** See `sim/misc/examples.c`

```

        .quad 0x008

# Matrix C

matC:
    .quad 0x000
    .quad 0x000
    .quad 0x000
    .quad 0x000

```

Your program(you are supposed to turn in) should consist of code that sets up a stack frame, invokes a function `matrix_and_xor`, and then halts. The function should be functionally equivalent to the C function `matrix_and_xor` shown in Figure Figure 1.

Note: To test your program you need to follow the below mentioned steps:

- a) In the `archlab-handout` directory you will find a directory named `matrix_test`. This directory has some sample matrices to be used as input for test cases and a script.
- b) To test your program you need to put just the function body of `matrix_and_xor` in a file named `matrix_and_xor.py` which is present in the directory `matrix_test`. Please do not copy anything except the function body. The initialization statements, call to `matrix_and_xor` function and stack initialization will be automatically taken care of. You are encouraged to have a look at `matrix_and_xor.py` in the `matrix_test` directory.

To execute the script enter the following command in the directory `matrix_test`

```
./test_matrix_and_xor.pl
```

- c) Before testing your function `matrix_and_xor` with the script you are advised to manually verify that your code is working. We will use a script similar to this one to grade. However, we will run your program for different inputs.
- d) This script assumes that you are passing the arguments in the same order as shown in the example c code in this handout i.e. `%rdi` takes size, `%rsi` takes matrix A, `%rdx` takes matrix B, and `%rcx` takes matrix C.

### **copy.py: Copy a source block to a destination block**

Write a program (`copy.py`) that copies a block of words from one part of memory to another (non-overlapping area) area of memory, computing the checksum (Xor) of all the words copied.

Your program should consist of code that sets up a stack frame, invokes a function `copy_block`, and then halts. The function should be functionally equivalent to the C function `copy_block` shown in Figure Figure 1. Test your program using the following three-element source and destination blocks:

```

.align 8
# Source block
src:
    .quad 0x00a

```

```

        .quad 0x0b0
        .quad 0xc00

# Destination block
dest:
        .quad 0x111
        .quad 0x222
        .quad 0x333

```

## 5 Part B

You will be working in directory `sim/seq` in this part.

Your task in Part B is to extend the SEQ processor to support the `iaddq`, described in Homework problems 4.51 and 4.52. To add this instructions, you will modify the file `seq-full.hcl`, which implements the version of SEQ described in the CS:APP3e textbook. In addition, it contains declarations of some constants that you will need for your solution.

### Building and Testing Your Solution

Once you have finished modifying the `seq-full.hcl` file, then you will need to build a new instance of the SEQ simulator (`ssim`) based on this HCL file, and then test it:

- *Building a new simulator.* You can use `make` to build a new SEQ simulator:

```
unix> make VERSION=full
```

This builds a version of `ssim` that uses the control logic you specified in `seq-full.hcl`. To save typing, you can assign `VERSION=full` in the Makefile.

- *Testing your solution on a simple Y86-64 program.* For your initial testing, we recommend running simple programs such as `asumi.yo` (testing `iaddq`) in TTY mode, comparing the results against the ISA simulation:

```
unix> ./ssim -t ../y86-code/asumi.yo
```

If the ISA test fails, then you should debug your implementation by single stepping the simulator in GUI mode:

```
unix> ./ssim -g ../y86-code/asumi.yo
```

- *Retesting your solution using the benchmark programs.* Once your simulator is able to correctly execute small programs, then you can automatically test it on the Y86-64 benchmark programs in `../y86-code`:

```
unix> (cd ../y86-code; make testssim)
```

This will run `ssim` on the benchmark programs and check for correctness by comparing the resulting processor state with the state from a high-level ISA simulation. Note that none of these programs test the added instructions. You are simply making sure that your solution did not inject errors for the original instructions. See file `../y86-code/README` file for more details.

- *Performing regression tests.* Once you can execute the benchmark programs correctly, then you should run the extensive set of regression tests in `../ptest`. To test everything except `iaddq` and leave:

```
unix> (cd ../ptest; make SIM=../seq/ssim)
```

To test your implementation of `iaddq`:

```
unix> (cd ../ptest; make SIM=../seq/ssim TFLAGS=-i)
```

For more information on the SEQ simulator refer to the handout *CS:APP3e Guide to Y86-64 Processor Simulators* (`simguide.pdf`).

## 6 Part C

You will be working in directory `sim/pipe` in this part.

The `ncopy` function in Figure 2 takes each element of `len`-element integer array `src`, copies it to a non-overlapping `dst`, and counts number of positive, negative and zero ints contained in `src`.

Figure 3 shows the baseline Y86-64 version of `ncopy`. The file `pipe-full.hcl` contains a copy of the HCL code for PIPE, along with a declaration of the constant value `IIADDQ`.

Your task in Part C is to modify `ncopy.y86` and `pipe-full.hcl` with the goal of making `ncopy.y86` run as fast as possible.

You will be handing in two files: `pipe-full.hcl` and `ncopy.y86`.

### Coding Rules

You are free to make any modifications you wish, with the following constraints:

- Your `ncopy.y86` function must work for arbitrary array sizes. You might be tempted to hardwire your solution for 64-element arrays by simply coding 64 copy instructions, but this would be a bad idea because we will be grading your solution based on its performance on arbitrary arrays.
- Your `ncopy.y86` function must run correctly with YIS. By correctly, we mean that it must correctly copy the `src` block *and* return (in `%rax`) the correct number of positives, (in `%rbx`) the correct number of negatives and (in `%rcx`) the correct number of zeros.
- The assembled version of your `ncopy` file must not be more than 500 bytes long. You can check the length of any program with the `ncopy` function embedded using the provided script `check-len.pl`:

```

1 /*
2  * ncopy - copy src to dst, and count number of positive,
3  * negative and zero ints contained in src array.
4  */
5 void ncopy(word_t *src, word_t *dst, word_t len)
6 {
7     word_t count_pos = 0, count_neg = 0, count_zero = 0;
8     word_t val;
9
10    while (len > 0) {
11        val = *src++;
12        *dst++ = val;
13        if (val > 0)
14            count_pos++;
15        else if (val < 0)
16            count_neg++;
17        else
18            count_zero++;
19        len--;
20    }
21 }

```

Figure 2: **C version of the ncopy function.** See `sim/pipe/ncopy.c`.

```
unix> ./check-len.pl < ncopy.yo
```

- Your `pipe-full.hcl` implementation must pass the regression tests in `../y86-code` and `../ptest` (without the `-i` flag that tests `iaddq`).

Other than that, you are free to implement the `iaddq` instruction if you think that will help. You may make any semantics preserving transformations to the `ncopy.yo` function, such as reordering instructions, replacing groups of instructions with single instructions, deleting some instructions, and adding other instructions. You may find it useful to read about loop unrolling in Section 5.8 of CS:APP3e.

## Building and Running Your Solution

In order to test your solution, you will need to build a driver program that calls your `ncopy` function. We have provided you with the `gen-driver.pl` program that generates a driver program for arbitrary sized input arrays. For example, typing

```
unix> make drivers
```

will construct the following two useful driver programs:



```

1 #####
2 # ncopy.ys - Copy a src block of len words to dst.
3 # Count the number of positive, negative and zero words contained in src.
4 #
5 # Include your name and ID here.
6 #
7 # Describe how and why you modified the baseline code.
8 #
9 #####
10 # Do not modify this portion
11 # Function prologue.
12 # %rdi = src, %rsi = dst, %rdx = len
13 ncopy:
14
15 #####
16 # You can modify this portion
17     # Loop header
18     xorq %rax,%rax           # count = 0;
19     andq %rdx,%rdx           # len <= 0?
20     jle Done                 # if so, goto Done:
21
22 Loop:  mrmovq (%rdi), %r10     # read val from src...
23         rmmovq %r10, (%rsi)    # ...and store it to dst
24         andq %r10, %r10        # val <= 0?
25         jle Npos              # if so, goto Npos:
26         irmovq $1, %r11
27         addq %r11, %rax        # Count positives in %rax - count_pos++
28         jmp Rest
29 Npos:   andq %r10, %r10        # Not positive
30         je Zero
31         irmovq $1, %r11
32         addq %r11, %rbx        # Count negatives in %rbx - count_neg++
33         jmp Rest
34 Zero:   irmovq $1, %r11
35         addq %r11, %rcx        # Count zeroes in %rcx - count_zero++
36 Rest:   irmovq $1, %r10
37         subq %r10, %rdx        # len--
38         irmovq $8, %r10
39         addq %r10, %rdi        # src++
40         addq %r10, %rsi        # dst++
41         andq %rdx,%rdx        # len > 0?
42         jg Loop              # if so, goto Loop:
43 #####
44 # Do not modify the following section of code
45 # Function epilogue.
46 Done:
47     ret
48 #####
49 # Keep the following label at the end of your function
50 End:

```

Figure 3: **Baseline Y86-64 version of the ncopy function.** See `sim/pipe/ncopy.ys`.

- `sdriver.yo`: A *small driver program* that tests an `ncopy` function on small arrays with 4 elements. If your solution is correct, then this program will halt with a value of 2 in register `%rax` after copying the `src` array.
- `ldriver.yo`: A *large driver program* that tests an `ncopy` function on larger arrays with 63 elements. If your solution is correct, then this program will halt with a value of 31 (0x1f) in register `%rax` after copying the `src` array.

Each time you modify your `ncopy.yo` program, you can rebuild the driver programs by typing

```
unix> make drivers
```

Each time you modify your `pipe-full.hcl` file, you can rebuild the simulator by typing

```
unix> make psim VERSION=full
```

If you want to rebuild the simulator and the driver programs, type

```
unix> make VERSION=full
```

To test your solution in GUI mode on a small 4-element array, type

```
unix> ./psim -g sdriver.yo
```

To test your solution on a larger 63-element array, type

```
unix> ./psim -g ldriver.yo
```

Once your simulator correctly runs your version of `ncopy.yo` on these two block lengths, you will want to perform the following additional tests:

- *Testing your driver files on the ISA simulator.* Make sure that your `ncopy.yo` function works properly with YIS:

```
unix> make drivers
unix> ../misc/yis sdriver.yo
```

- *Testing your code on a range of block lengths with the ISA simulator.* The Perl script `correctness.pl` generates driver files with block lengths from 0 up to some limit (default 65), plus some larger sizes. It simulates them (by default with YIS), and checks the results. It generates a report showing the status for each block length:

```
unix> ./correctness.pl
```

This script generates test programs where the result count varies randomly from one run to another, and so it provides a more stringent test than the standard drivers.

If you get incorrect results for some length  $K$ , you can generate a driver file for that length that includes checking code, and where the result varies randomly:

```

unix> ./gen-driver.pl -f ncopy.ys -n K -rc > driver.ys
unix> make driver.yo
unix> ../misc/yis driver.yo

```

The program will end with register `%rax` having the following value:

**0xaaaa** : All tests pass.

**0xbbbb** : Incorrect count

**0xcccc** : Function `ncopy` is more than 500 bytes long.

**0xdddd** : Some of the source data was not copied to its destination.

**0xeeee** : Some word just before or just after the destination region was corrupted.

- *Testing your pipeline simulator on the benchmark programs.* Once your simulator is able to correctly execute `sdriver.ys` and `ldriver.ys`, you should test it against the Y86-64 benchmark programs in `../y86-code`:

```

unix> (cd ../y86-code; make testpsim)

```

This will run `psim` on the benchmark programs and compare results with `YIS`.

- *Testing your pipeline simulator with extensive regression tests.* Once you can execute the benchmark programs correctly, then you should check it with the regression tests in `../ptest`. For example, if your solution implements the `iaddq` instruction, then

```

unix> (cd ../ptest; make SIM=../pipe/psim TFLAGS=-i)

```

- *Testing your code on a range of block lengths with the pipeline simulator.* Finally, you can run the same code tests on the pipeline simulator that you did earlier with the ISA simulator

```

unix> ./correctness.pl -p

```

## 7 Evaluation

The lab is worth 100 points: 30 points for Part A, 30 points for Part B, and 40 points for Part C.

### Part A

Part A is worth 30 points, 10 points for each Y86-64 solution program. Each solution program will be evaluated for correctness, including proper handling of the stack and registers, as well as functional equivalence with the example C functions in `examples.c`.

The program `max.ys` will be considered correct if your code works correctly for every possible input, and its respective `max_list` functions returns the max `0xc00` in register `%rax`.

The program `matrix_and_xor.ys` will be considered correct if your code works correctly for every possible input, and the destination matrix and register `%rax` contain the correct values

The program `copy.y8` will be considered correct if your code works correctly for every possible input, and the `copy_block` function returns the sum `0xcba` in register `%rax`, copies the three 64-bit values `0x00a`, `0x0b`, and `0xc` to the 24 bytes beginning at address `dest`, and does not corrupt other memory locations.

## Part B

This part of the lab is worth 30 points:

- 15 points for passing the benchmark regression tests in `y86-code`, to verify that your simulator still correctly executes the benchmark suite.
- 15 points for passing the regression tests in `pctest` for `iaddq`.

## Part C

This part of the Lab is worth 40 points: **You will not receive any credit if either your code for `ncopy.y8` or your modified simulator fails any of the tests described earlier.**

- All the points are for performance. To receive credit here, your solution must be correct, as defined earlier. That is, `ncopy` runs correctly with `YIS`, and `pipe-full.hcl` passes all tests in `y86-code` and `pctest`.

We will express the performance of your function in units of *cycles per element* (CPE). That is, if the simulated code requires  $C$  cycles to copy a block of  $N$  elements, then the CPE is  $C/N$ . The PIPE simulator displays the total number of cycles required to complete the program. The baseline version of the `ncopy` function running on the standard PIPE simulator with a large 63-element array has an average CPE of 18.30.

Since some cycles are used to set up the call to `ncopy` and to set up the loop within `ncopy`, you will find that you will get different values of the CPE for different block lengths (generally the CPE will drop as  $N$  increases). We will therefore evaluate the performance of your function by computing the average of the CPEs for blocks ranging from 1 to 64 elements. You can use the Perl script `benchmark.pl` in the `pipe` directory to run simulations of your `ncopy.y8` code over a range of block lengths and compute the average CPE. Simply run the command

```
unix> ./benchmark.pl
```

to see what happens. For example, the baseline version of the `ncopy` function has CPE values ranging between 36.00 and 18.27, with an average of 19.30. Note that this Perl script does not check for the correctness of the answer. Use the script `correctness.pl` for this.

You should be able to achieve an average CPE of less than 12.65. If your average CPE is  $c$ , then your score  $S$  for this portion of the lab will be:

$$S = \begin{cases} 0, & c > 12.65 \\ 20 \cdot (12.65 - c), & 10.65 \leq c \leq 12.65 \\ 40, & c < 10.65 \end{cases}$$

By default, `benchmark.pl` and `correctness.pl` compile and test `ncopy.js`. Use the `-f` argument to specify a different file name. The `-h` flag gives a complete list of the command line arguments.

## 8 Handin Instructions

- You will be handing in three sets of files:
  - Part A: `max.js`, `matrix_and_xor.js`, and `copy.js`.
  - Part B: `seq-full.hcl`.
  - Part C: `ncopy.js` and `pipe-full.hcl`.
- Put your files for each part in a separate directory `part_a`, `part_b`, `part_c`.
- Compress all files into a single `.zip` file and submit it in the Moodle link for Architecture Lab.

## 9 Hints

- By design, both `sdriver.yo` and `ldriver.yo` are small enough to debug with in GUI mode. We find it easiest to debug in GUI mode, and suggest that you use it.
- If working on remote machines, to have the GUI version of `ssim` and `psim`, you can try FastX, which can be found at this link: [help.cselabs.umn.edu/offsite/fastx](http://help.cselabs.umn.edu/offsite/fastx). Make sure that you are connected to some secure network viz, UofM secure and not an open network like UofM guest. This utility lets you have the GUI for the remote machine.
- Alternatively, If you are running in GUI mode on a Unix server, make sure that you have initialized the `DISPLAY` environment variable:

```
unix> setenv DISPLAY myhost.edu:0
```

- With some X servers, the “Program Code” window begins life as a closed icon when you run `psim` or `ssim` in GUI mode. Simply click on the icon to expand the window.
- With some Microsoft Windows-based X servers, the “Memory Contents” window will not automatically resize itself. You’ll need to resize the window by hand.
- The `psim` and `ssim` simulators terminate with a segmentation fault if you ask them to execute a file that is not a valid Y86-64 object file.